

## MapReduce

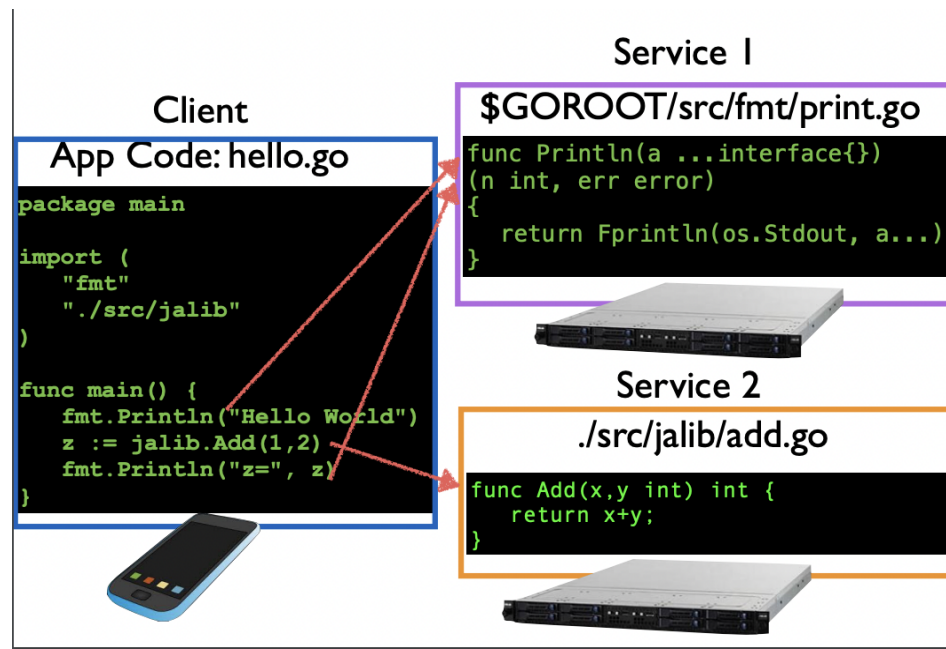
### 1. Example

- a. File with #pid, age, diag, glucose level, heart rate, ...  
135, 36, d1, 120, ..., ....
- b. Problem: Average glucose level for all patient between 40 and 60 years old and group by diagnosis code in terms of MapReduce
  - i. Map phase
    1. input: partition of file
    2. Output (require intermediate key value pair)
      - a. Key: diag
      - b. value: glucose level
    3. The filtering of age (age between 40 and 60) within the map phase
  - ii. Data Shuffle (between map and reduce phase)
    1. The distribution of key-value pair inside map to appropriate reduce function
  - iii. Reduce phase
    1. Reduce (key, iterator)  $\rightarrow$  reduce (diag, [glucose 1, glucose 2, ...])
    2. Diag  $\rightarrow$  iterator
    3. Sum the glucose level per diagnosis
    4. Record the count for each key
    5. Divide Sum by the count to calculate the average
    6. Output  $\rightarrow$  (diagnosis, average)  $\rightarrow$  [(d1, avg), (d2, avg), ...]
- c. What happens if there exists a skew in data (d1 has huge amounts of data while other diagnoses do not contain many data)?
  - i. Ex) (d1, g1), (d1, g2), (d2, g3), (d3, g4), (d1, g5)  $\rightarrow$  many diagnosis for d1
  - ii. Combiner function between map phase and reduce phase that adds all the glucose values (g) for each key diagnosis (d)  $\rightarrow$  reduce intermediate amount of data
    1. d1, (322, 3)
    2. d2, (g3, 1)
    3. d3, (g4, 1)

### 2. RPC - Remote Procedure Calls

- a. One of the bread and butter building blocks for distributed system construction
- b. Hopefully a particular RPC infrastructure is boring once you get the basic idea and have read the docs
- c. Our goal today is to both get a handle on the idea and use

### 3. The idea: libs as services



a.

b. Functions as services:

- i. Functions do not have to be stored in the same machine
- ii. Different applications/users can use the same function

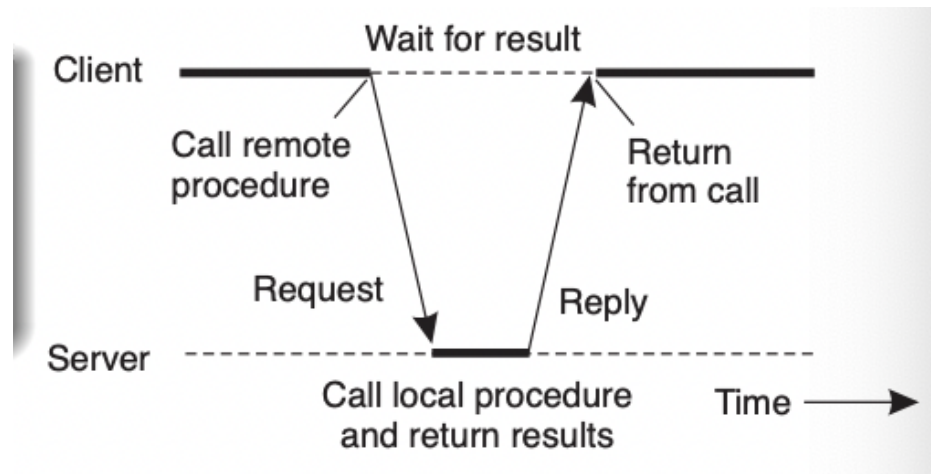
### 4. Remote Procedure Call (RPC)

a. Observations:

- i. Application developers are familiar with simple procedure model
- ii. Well-engineered procedures operate in isolation (black box)
- iii. There is no fundamental reason not to execute procedures on separate machine


b. Conclusion

- i. Communication between caller & callee can be hidden by using procedure-call mechanism



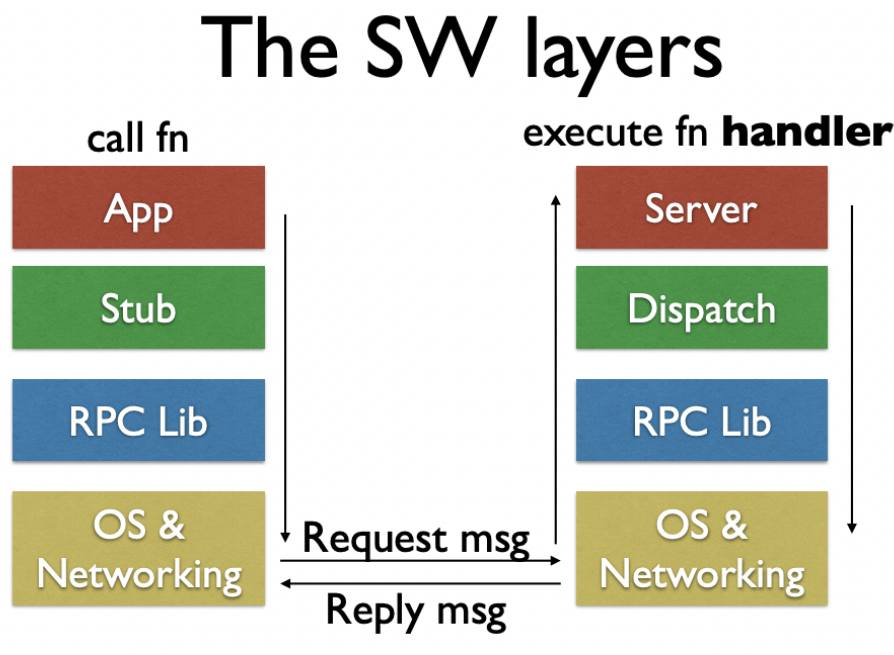
ii.

## 5. Interactions/Msgs



a.

## 6. The SW layers



a.

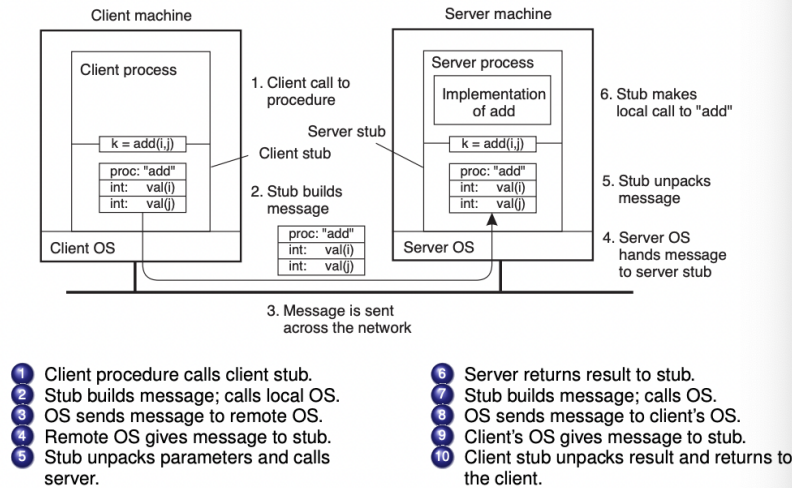
b. Program initiates actual RPC

c. The application using RPC communicates and request the RPC

d. How to do this: take an object → break into bytes → send to RPC library → send to operating system → send it to other server and repeat the process (reconstruct the object → ... → send to application on the server side)

## 7. Under the covers


# Under the covers



a.


## 8. Threads

■ ■ ■



a.

- b. Might have multiple threads all of which could concurrently be making rpc's to one more servers




- c.
- d. Rpc handlers might take a long time - often use threads to execute many rpcs concurrently (thread per handler execution)

#### 9. Failures?

- a. Lost packet
- b. Broken network
- c. Slow server
- d. Crashed server
- e. From the client's perspective, failures typically mean that client is waiting for a reply that will never come

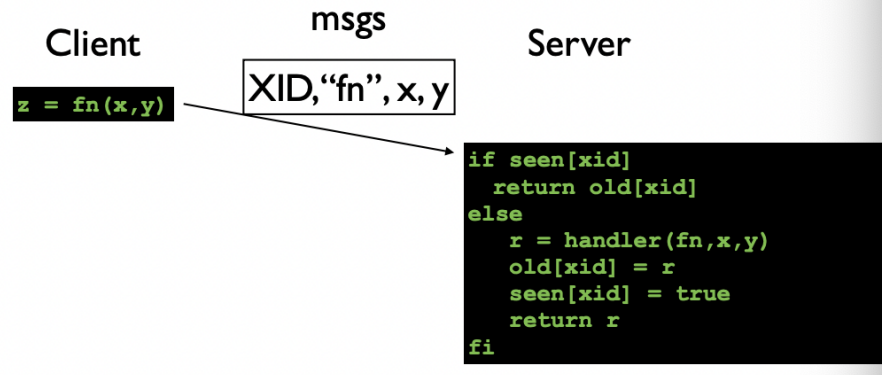
#### 10. If (no reply in X seconds) then?

- a. At least once failure model



- i.
- ii. Regardless of failures execute the rpc at least once
- iii. While true{
  - Send request
  - Wait x seconds for reply
  - If reply return

- iv. As long as eventually some or something fixes the problem (eg. robot server, fix network), then this will always work
- v. It does work for read-only operations
- vi. Or you have a strategy for duplicates (which later labs will require)
- b. At most once
  - i. Server detects duplicates and not execute handler
  - ii. Easy way to detect duplicates
    - 1. RPC id



a.

```

if seen[xid]
  return old[xid]
else
  r = handler(fn,x,y)
  old[xid] = r
  seen[xid] = true
  return r
fi
  
```

b.

- c. Introduce a unique id per-RPC invocation and some storage
  - if time runs out, and tries again, it uses the same id as the old id
- d. Server side: Server has to maintain all these ids (which can lead to running out of memory)
- e. This works but there are some issues
- f. How do we delete things from old and seen?
  - i. Get an ack from the client for XID for which it has received responses