


Google File System (GFS)

1. Why does Google need it?
 - a. Handle huge amount of data and failures
 - b. GFS → one of first design to exchange performance and consistency
2. MapReduce
 - a. Storage is provided by inexpensive IDE disks attached directly to individual machines. A distributed file system developed in-house is used to manage the data stored on these disks. The file system uses replication to provide availability and reliability on top of unreliable hardware
 - b. The more web page point to other webpages, the more popular (collect link among webpages to create a graph where nodes are webpages)
3. Topic: scalability, fault-tolerance, and data consistency
4. Consistency
 - a. A correctness condition
 - b. Inconsistent → read/write the data at the same time
 - c. Important when data is replicated and concurrently accessed by applications
 - i. If a write is performed, what value then will later read observe?
 - ii. What if two writes are performed concurrently?
 - iii. What if the reader is different from the writers?
 - d. Data's consistency model defines what is correct




- e.
- f. Reason why you store multiple replication of same data → data management (split the work), ensure that server is available even after one replication is down (fault tolerance)

- g. Strong consistency
 - i. Read always returns the value of the most recent write
 - h. Weak consistency
 - i. Reads may return “stale” data (not the value of the most recent write)
 - ii. Implement weak consistency due to performance
 - iii. The benefits we get from using weak consistency is great
 - i. In a weak consistency model, inconsistency is ... OK!
5. Long History of Consistency
- a. Decades of development from architecture, operating systems, database, and distributed communities:
 - i. Concurrent processors with private caches accessing a shared memory
 - ii. Concurrent clients accessing a file system
 - iii. Concurrent transactions on distributed database
6. Consistency Today?
- a. Important data maybe replicated to prevent it from being lost due to a failure (durability)
 - b. Important services maybe replicated to handle increasing demand, or guarantee availability in the case of a failure
 - c. Any mutations of a replicated value has the potential to introducing inconsistencies


7. Sources of inconsistency

- a. Multiple threads, multiple clients



b.

- c. Concurrency!

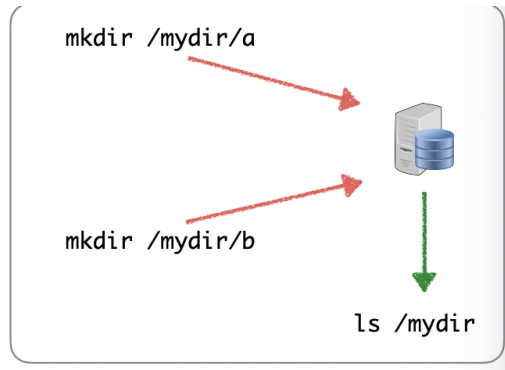


d.

- e. Network failure or machine failure occurs

8. “Ideal” correctness guarantees

- a. In ideal world, replicated system behaves like a non-replicated system



b.

- i. “Single copy” serialized access
- ii. After a write, all reads will always see update value
- iii. Concurrent writes to same file independently done in a well defined order