



## Raft

1. Two general problems
  - a. Two generals need to coordinate an attack
  - b. Must agree on the time to attack
  - c. They win only if they attack simultaneously
  - d. Communicate through messengers
  - e. Messengers may be killed on their way
2. Solve the two generals problem
  - a. G1 sends time of attack to G2
  - b. Problem: how to ensure G2 received message?
    - i. Solution: let G2 ack receipt of msg
  - c. Problem: how to ensure G1 received ack
    - i. Solution: let G1 ack the receipt of the ack
  - d. This problem is impossible to solve
3. Two general problem
  - a. Applicability to distributed systems
  - b. Two nodes need to agree on a value
  - c. Communicate by messages using an unreliable channel
  - d. Consensus is a core problem
4. Raft
  - a. Fault tolerance through replicated state machines (RSM)
    - i. Fault tolerance
      1. Execution (reassign the task to other workers) → in MapReduce
      2. Primary backup
      3. Etc.
  - b. RAFT: much more complete design than the PAXOS paper
  - c. Fundamentally, Raft bakes in notion of a leader and log
5. Big picture
  - a. Have bunch of clients and servers
  - b. Servers are accumulating logs of client commands




c.

- d. Goal is to keep the logs in the servers identical (when given failures → server and network)



- e.
- f. At any given time, there is at most one leader (from the server) → example Server1
- g. Other server (server2 and server3) become followers here



- h.
- i. Clients are supposed to send requests to the Leader
- j. Client0 sends request A to leader (server1)
- k. Leader updates the command to its log when it receives request A from client
- l. Sends concurrent request to followers (server2 and server3)
- m. Once the leader believes majority of the followers (including itself) have command in there log it announces commitment and that command can be executed safely → send back result to the client

## 6. What matters

- NOT: is it more or less understandable that PAXOS
- NOT: is it better, faster, more efficient
- RATHER: it is a relatively complete tutorial on how to build a RSM system


## 7. PAXOS - lots of messages

- Raft can replicate whole sequence of commands while PAXOS is not (reduces the communication between the protocols greatly)
- Whole bunch of independent PAXOS instances - too slow & requires lots of messages → in practice optimize so that common case does not need this


## 8. What are we hoping for?

- Tolerate failure of a minority of followers (if majority fails, RAFT cannot do anything about it)
- Convergence to one version of the log


- c. No undos of logs
  - d. Only execute client requests (and thus reply to client) when it is “committed” – can’t go back and unexecuted or un-tell client a response
9. No Leader failure
- a. But: Followers’ crash, Network loses messages



- b.
- c. LOGS same by force:
  - i. Leader
    - 1. Appends to its log
    - 2. In parallel Sends index (sequence number that is command in the log → 1,2,3,4,etc in log) and “term” (sequence number that indicates a phase in the RAFT protocol → each time there is a new election, the term increases) of prior entry as part of AppendEntries RPC (concurrent requests)
  - ii. Followers
    - 1. If prior index and term not in log REJECT if match then log is consistent
- d. Here, S3 will not append the command while S2 will append the command
- e. LOGS same by force: If a follower is found inconsistent then leader rolls back follower and then rolls it forward



- f.
- g. Leader does not care what is in your log; it forces its view on all followers! (when leader sends the current command, it also sends the previous index and term number and the previous command → if the command is not the same, the follower replies back to the leader that the numbers/commands do not match → leader sends the previous command/numbers to that follower so that the follower can have the same command/number)



- h.
- i. When follower rejects the command, it sends its previous command/numbers to the leader → leader jumps to that command/numbers


#### 10. 2 common divergence

- a. Follower crashed for a while (disconnected)
  - i. The leader simply rolls follower forward
- b. A prior leader crashed while sending out some AppendEntries messages
  - i. Then leader both rolls follower back and forward
- c. Consistency through leader driven reconciliation

#### 11. Summary


- a. Tolerate failure of a minority of followers
  - i. Various times states “wait for majority”, “when majority” (majority matters → if majority fails, then we are done)
- b. Convergence on one version of the log
  - i. By reconciliation to leader’s log
- c. Only execute “committed” requests
  - i. Not a problem if leader does not crash and waits for a majority to respond that append is durable - then sends out a commit point (piggy backed or AppendEntries RPC)

#### 12. HeartBeat to followers



- a.
- b. appendEntries done repeatedly by leader to all followers (null if necessary)
- c. Followers initiate an “Election” if heartbeat timeout expires
- d. The purpose is for the followers to acknowledge that there is a leader
- e. Each server has a timeout → followers immediately initiate a new leader election

### 13. Time broken into Terms



- a.
  - b. Detect that a leader has “crashed”
  - c. Terms are distinct phases of the protocol
  - d. Start an election to initiate a new term with a new leader (term cannot go backwards)
  - e. New leader has to pick up the pieces (when there is no leader, in the outside world, the system appears dead → cannot take any request)
  - f. Remember a leader can crash in the middle of anything
14. Leader election
- a. Fewer than 2 leaders any time (0 and 1 is ok)
  - b. Ideally greater than 0 (so the most ideal is one leader)
  - c. Raft has separate mechanisms for these
15. Fewer than 2 (at most one) leader elected per term
- a. When a server gets a vote request it is allowed to vote yes or no BUT it is only allowed One vote in any given term - it can only say yes to one candidate
  - b. Candidate must collect yes votes from a majority to become the leader
  - c. Therefore, only one candidate become the leader for a term
16. Reason to use the majority
- a. Only requiring majority (not everyone) means that you can tolerate failure of a minority!
  - b. Avoids split brain
  - c. Any two majorities must overlap with one server - prior majority must share a server with current majority - ensures continuity
  - d. New leader will be sure to know about all previous decisions (committed entries)