


## Spark

### 1. Why Spark?

- General-purpose cluster computing framework
- Better performance than Hadoop (open-source MapReduce)
- Open source protocol
- User-friendly fault-tolerant abstractions (RDDs)
- Rich dataflow API (incl. `join()`, `flatMap()`, `groupByKey()`, etc.)
- A big commercial success

### 2. MapReduce

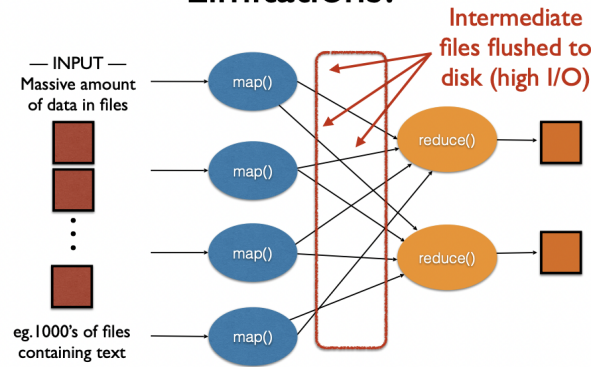


a.

- User must specify two functions → map and reduce functions

### 3. Limitations of MapReduce


#### Limitations:



a.

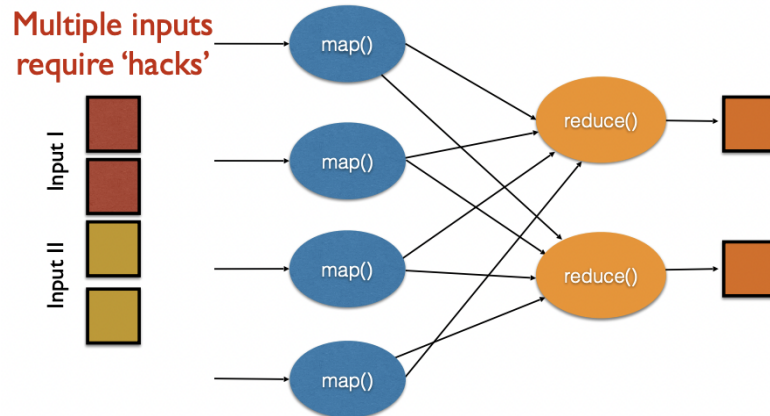
- Intermediate files flushed to disk (high I/O)
- Users do not know anything about the file regarding memory and etc. but is up to the system

- d. Iterations require manual work by users → can only implement one map and one reduce function at a time (does not support building map and reduce functions together) → users have to know where the output data is stored to send to another mapReduce protocol repeatedly



e.

- f. Multiple inputs require ‘hacks’



g.

- h. MapReduce does not support more than one inputs (only considers that we have a big logical file separated into many text files that belong to one input data)
- i. If you want to insert multiple inputs, you somehow have to include this information into the input record
- j. Data is stored on disk
- k. How would map() and reduce() distinguish between k-v pairs from different inputs?
- Put it inside the record manually by the user
- l. Recognizing this problem, researchers have developed specialized frameworks for some applications that require data reuse
- m. They do not provide abstractions for more general reuse, e.g., to let a user load several datasets into memory and run ad-hoc queries across them


#### 4. Spark

- a. A new abstraction called resilient distributed datasets (RDDs) that enables efficient data reuse in a broad range of applications. RDDs are fault-tolerant, parallel data structures that let users explicitly persist intermediate results in memory, control their partitioning to optimize data placement, and manipulate them using a rich set of operators.
- b. RDD
  - i. Abstraction
  - ii. Data structure that users can manipulate
  - iii. Fault-tolerant
  - iv. Collection of data records that is physically distributed across machines

#### 5. Dataflow Graphs




- a.
- b. Each data takes inputs as RDDs (one or more) and generates RDDs
- c. Example of wordcount operation by using Spark




```
val textFile = sc.textFile("hdfs://...")
val counts = textFile.flatMap(line => line.split(" "))
                      .map(word => (word, 1))
                      .reduceByKey(_ + _)
counts.saveAsTextFile("hdfs://...")
```

- d.
- e. Just specify the logic
- f. textFile → name for RDD (handle the file you read from disk)
- g. flatMap takes each record in the RDD (as line) and splits it into words
- h. The output (input of words) of it goes as input to map
- i. Map creates a new record and outputs RDD that is given input to ReduceByKey
- j. The ReduceByKey outputs a RDD
- k. Counts → another RDD that defines the output of the operator
- l. The last line of code saves the RDD in the distributed file system
- m. Everything in Spark is made of RDDs → input, output, between the functions


## 6. Spark Application



- a.
  - b. User writes application that is managed by the 'driver' (just like coordinator in MapReduce)
  - c. Driver defines one or more RDD and invokes actions on them
  - d. Workers can store data in RAM across operations (unlike MapReduce)
  - e. More like writing a program than MapReduce - explicitly naming data and operating on it
7. RDD (collection of data records distributed in machines → memory or disk)
- a. RDDs provide an interface based on coarse-grained transformations

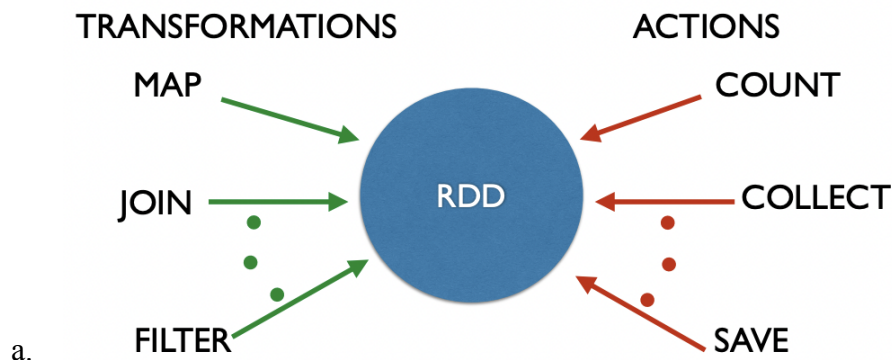


- b.
  - c. Generated by an input dataset or by existing RDD
  - d. Coarse grain operations that hide fine-grain memory R/Ws
  - e. Same operation across many items within the object
  - f. Restricted interface provides basis for Fault-Tolerance - log of operations per RDD (lineage)
8. Read-Only & SIMD (ensures fault tolerance)



- a.
- b. Created never modified
- c. Computations is a series of transformations that produce create a new RDD
- d. Partitions of data items
- e. Data parallel model of computing - single operation applied in parallel to many data items


## 9. Lazy Evaluation



- a.
- b. Green
  - i. Compute a new RDD from existing RDDs
  - ii. This just specifies a plan
  - iii. Runtime is lazy - doesn't have to materialize (compute) so it doesn't
    - 1. Does not actually process data available until it has to
    - 2. No RDD is generated, it is stored in driver
    - 3. The output does not collect all the RDDs, just the output
    - 4. Action → specific
- c. Red
  - i. Where some effect is requested: result to be stored, get specific value, etc. causes RDDs to materialize - return a value to app or to stable storage

## 10. Persistence


- a. Users can indicate which RDD's they will reuse and choose a storage strategy for them (eg. in-memory)



- b.
- c. Users can set priorities of the output (ex. Written to disk first when there is no enough memory)

## 11. Partitioning

- Users can ask that an RDD's records be partition across machine based on a key in each record



- Output has all pairs of records of the city

## 12. Logistic Regression

```
val points = spark.textFile(...)
                        .map(parsePoint).persist()
var w = // random initial vector
for (i <- 1 to ITERATIONS) {
  val gradient = points.map{ p =>
    p.x * (1/(1+exp(-p.y*(w dot p.x)))-1)*p.y
  }.reduce((a,b) => a+b)
  w -= gradient
}
```

*transformations* (pointing to `.map(parsePoint)`)

*action* (pointing to `.reduce((a,b) => a+b)`)

*(cf. Table 2 in the paper)*

- `w` is sent with the closure to the nodes
- Points → RDD (keeping memory), `persist()` materializes in memory by default assuming memory is enough
- Map operators → data transformations (iterate through all record and apply the action (reduce))
- Action → materialize the RDD
- `w` is sent with the closure to the nodes

```
val points = spark.textFile(...)
                        .map(parsePoint).persist()
var w = // random initial vector
for (i <- 1 to ITERATIONS) {
  val gradient = points.map{ p =>
    p.x * (1/(1+exp(-p.y*(w dot p.x)))-1)*p.y
  }.reduce((a,b) => a+b)
  w -= gradient
}
```

*for each point p in points do:*  
*p.x \* f(w,p) \* p.y*

f.

- g. For each point  $p$  in points do:
  - i.  $P.x * f(w,p) * p.y \rightarrow$  not materialized
- h. Reducer sums all input vectors and materialize a new RDD 'gradient'
- i.