

Spark (cont.)

1. Overview

- a. RDD → represents dataset (collection of records that is physically distributed in class of machines)
- b. RDDs → support operations
 - i. Passed into data operator (as inputs and outputs)
 - ii. RDDs are mutable data structures → generates fault tolerance
- c. Spark & MapReduce
 - i. Spark supports multiple inputs
 - ii. Spark gives more control on intermediate data (specify as the user to keep the data on disk or memory unlike MapReduce), MapReduce stores data in local disk on mapper
 - iii. Programming model is different
 - 1. MapReduce → define map and reduce functions (singular) and may have to implement many different versions
 - 2. Spark → more intuitive



- d.
- e. All the nodes are data file operators
- f. Data file operator → each instance applies the logic to the joint partitions
- g. Task parallelism → applying different task on the same data
 - i. ex) Multi Tread → different tasks on the same data (concurrent task)
- h. Data parallelism → same function, different data
- i. In reality, system breaks down into series of tasks that look like Map and reduce functions
- j. Workers always communicate with the driver/master (so that it can give the output of one worker to another worker as input)
- k. Transformations (map, join, filter, reduce by key) → take as inputs as RDD and produce RDDs and pass down the output to the next operator in the graph
- l. Pipeline allows data to flow and make data progress in parallel
 - i. Makes better use of resources
- m. Do we have pipeline in map and reduce?
 - i. No, we wait for the mappers to finish and start the reduce phase

- n. Transactions/actions → transformations pipeline the output to the next operator and actions output/materialize

```
val points = spark.textFile(...)
                    .map(parsePoint).persist()
var w = // random initial vector
for (i <- 1 to ITERATIONS) {
  val gradient = points.map{ p =>
    p.x * (1/(1+exp(-p.y*(w dot p.x)))-1)*p.y
  }.reduce((a,b) => a+b)
  w -= gradient
}
```

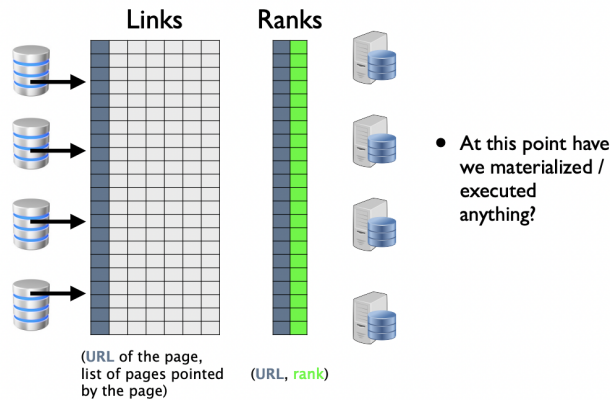
- o.
- i. Point → RDD
 - ii. map → transformation
 - iii. Gradient → RDD
 - iv. For loop → Anonymous function that defines logic of the map
 - v. Reduce → action that materializes gradient (because we need to subtract gradient)

2. PageRank

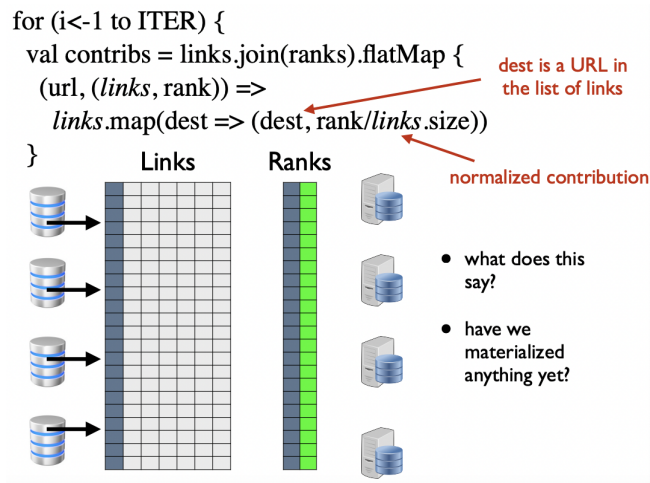
```
val links = spark.textFile(...).map(...).persist()
var ranks = // RDD of (URL, rank) pairs
for (i <- 1 to ITERATIONS) {
  // Build an RDD of (targetURL, float) pairs
  // with the contributions sent by each page
  val contribs = links.join(ranks).flatMap {
    (url, (links, rank)) =>
      links.map(dest => (dest, rank/links.size))
  }
  // Sum contributions by URL and get new ranks
  ranks = contribs.reduceByKey((x,y) => x+y)
                    .mapValues(sum => a/N + (1-a)*sum)
}
```

- a.
- b. Links, ranks → RDDs
 - c. links.join(ranks) → graph with two inputs (links and ranks) that go as input in join(transformation) that generates a new RDD that goes input into flatMap, which also produces an RDD
 - d. flatMap will apply the logic inside and will generate the records
 - e. Input format of flatMap → (url, and tuple of (links, rank)), and therefore, the links.join(ranks) should return such format
 - f. For each record in the collection, apply logic (desk → (dest, rank/links.size))

```
val links = spark.textFile(...).map(...).persist()
var ranks = // RDD of (URL, rank) pairs
```

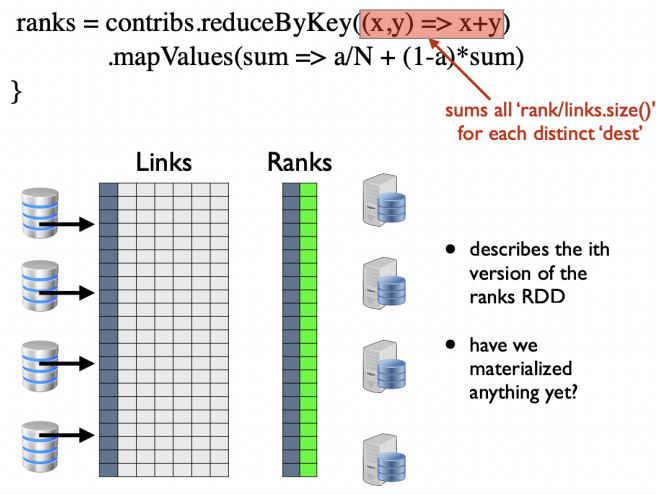


g.



h.

i. Join will join the two tables on the common attribute (URL) and will produce a RDD



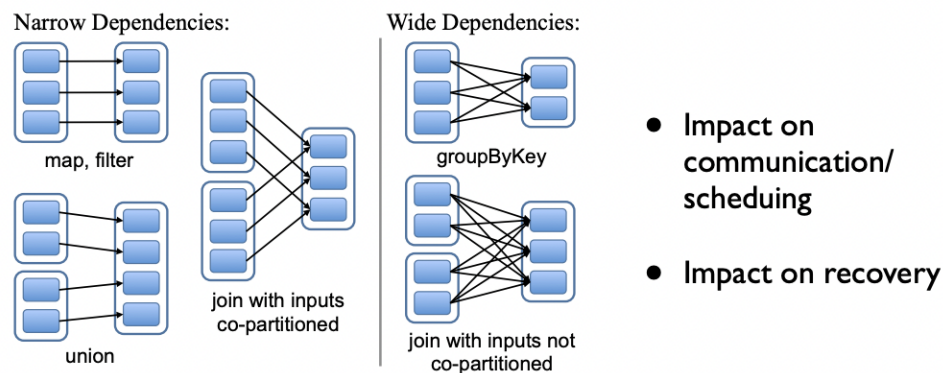
j.

k. Output of the flatmap is inserted as input to reduceByKey (transformation)

3. What an RDD is composed of

Operation	Meaning
<code>partitions()</code>	Return a list of Partition objects
<code>preferredLocations(<i>p</i>)</code>	List nodes where partition <i>p</i> can be accessed faster due to data locality
<code>dependencies()</code>	Return a list of dependencies
<code>iterator(<i>p</i>, <i>parentIters</i>)</code>	Compute the elements of partition <i>p</i> given iterators for its parent partitions
<code>partitioner()</code>	Return metadata specifying whether the RDD is hash/range partitioned

-
- RDD dependencies → which RDD depends on which RDD
- This means that if something fails, we know the RDDs that depend on, so we can handle fault tolerance
- Dependencies store parent RDDs, and data in case of failure in order to reconstruct the partition



-
-
-
- Why does an RDD carry metadata on its partitioning?
 - Each RDD also has some information about the way it is partitioned
 - Transformations that depend on multiple RDDs know whether they need to shuffle data (wide dependency) or not (narrow) and allows users control locality and reduce shuffles
- Handling Faults
 - No replication by default - even when persisted (because it is expensive)
 - Ram or disk node failure could mean permanent loss of data
 - Assume heartbeats used to detect lost worker

- i. If worker lost need a way to recompute it's partitions - will need all dependent partitions
 - ii. Recompute if they can't be found in RAM or on disk
 - iii. If wide dependency will need all partitions of dependent RDD if narrow then only one partition
- c. So two mechanisms enable recovery from faults: lineage and policy of what partitions to persist (driver and actual replication and partition strategy)