

## Data for Web Apps

1. Data and the web
  - a. Nearly every web application is backed by a database of some sort
  - b. Even static pages might have data requirements
  - c. Web-based languages have become optimized for aggregating data into HTML
  - d. It isn't just dynamic data that we are interested in...
2. Three kinds of data
  - a. Generally speaking we can group data into three buckets
    - i. User data: Things that are specific to each user, such as locale, preferences...
    - ii. Common data / content: Items that are present for all users
    - iii. Infrastructure data: Data required to operate the site
3. Four data activities
  - a. We often refer to data operations as implementing CRUD:
    - i. Create
    - ii. Read
    - iii. Update
    - iv. Delete
  - b. These operations (with the possible exception of Read) need to be atomic
4. Atomicity
  - a. We almost never build an app that has only one user
  - b. We need to account for CRUD operations in which several users might need to access the same data
  - c. The problem is that in modern operating systems, processes and threads are constantly switching, often in mid-instruction
  - d. Further, a data operation might take several steps...this is called a transaction
  - e. How can we guarantee that once a process has started to update data in a table, it will complete before another process starts to update the same data?
5. Locking
  - a. Most databases implement some sort of locking
  - b. The idea is that the db is locked when one thread or process is working on a dataset, preventing other threads from changing the data midstream
  - c. Locks can be either at the table level or row level
  - d. It's important to know which it is...it makes a difference in how you write db access code on the app side

## 6. ACID

- a. In most cases data encompasses information that we want to keep track of
- b. When we think of transactions in a database, we can implement a set of properties formalized by Jim Gray back in the 70s
- c. We extend the notion of atomicity...
  - i. A: Atomicity — all-or-nothing transactions
  - ii. C: Consistency — any transaction will result in the database being in a valid state
  - iii. I: Isolation — if transactions are done concurrently, the result is the same state that would have been reached had the transactions been done serially
  - iv. D: Durability — when a transaction has been committed to the database, it will remain in the dataset until it is updated by another transaction, even if the power is lost
- d. This applies primarily to traditional data sources such as relational databases
- e. We'll see that it isn't always strictly necessary to implement ACID transactions

## 7. Traditional data sources: mainframes

- a. Mainframes such as IBM's OS/360 series and z/OS machines are optimized for data
- b. Even though we tend to think of mainframes as old-school, there are still a huge number deployed
- c. We typically see these in financial services, insurance, health care, and related industries
- d. The applications that sit on top of these data sources are tightly coupled and custom-built for each client

## 8. Fixed-width (or fixed-length) data

- a. It's common to find data sources that expect data to be in fixed-width format
- b. This is common in mainframe applications (AS/400 etc)
- c. The format originated with punch cards
- d. In fixed-width, information fits into a certain number of characters
- e. For example, the first 5 characters might indicate the record type, the next 5 characters are a transaction ID, the next 4 a key, and so on

## 9. Delimited records (EDI)

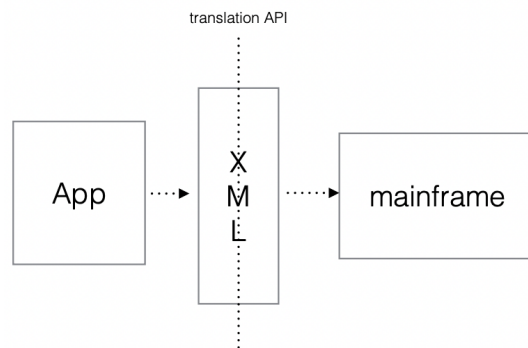
- a. EDI is Electronic Data Interchange
- b. Many industries move data back and forth in fixed-format message
- c. There are several EDI standards (X12, EDIFACT, and so on)
- d. Healthcare, government, finance, transportation / shipping, supply chain all use EDI
- e. Each field in the record is delimited by one or more special characters

## 10. EDIFACT sample message

```
UNA:+.? '
UNB+IATB:1+6XPPC+LHPPC+940101:0950+1 '
UNH+1+PAORES:93:1:IA '
MSG+1:45 '
IFT+3+XYZCOMPANY AVAILABILITY '
ERC+A7V:1:AMD '
IFT+3+NO MORE FLIGHTS '
ODI '
TVL+240493:1000::1220+FRA+JFK+DL+400+C '
PDI++C:3+Y::3+F::1 '
APD+74C:0:::6++++++6X '
TVL+240493:1740::2030+JFK+MIA+DL+081+C '
PDI++C:4 '
APD+EM2:0:1630::6++++++DA '
UNT+13+1 '
a. UNZ+1+1 '
```

## 11. Parsing EDI

- Typically a commercial library will be used to parse EDI messages into objects or other data formats
- Reading and writing EDI messages conforms to strict error- and formatchecking
- It isn't unusual for a batch of EDI messages to be FTPd to a destination rather than be sent in real time
- However!
- There's increasing demand to provide app- and web- based access to this data
- It's not cost effective to simply replace the mainframe...there's a large investment and they actually are extremely good at processing data
- The solution is usually middleware that sits between the client application and the mainframe



- 
- 
- 
- 
- 
- 
- 
- 
- From the app's point of view, all data operations are accomplished via XML requests and response
- This is a pretty common pattern
- It allows us to aggregate mainframe data operations with other data sources...i.e. the app requests user accounts data, and the account information comes from a traditional database but the account balances come from the mainframe

## 12. Traditional data sources: RDB

- An RDB is a relational database such as Oracle, DB2, or MySQL
- In RDBs data is stored in tables
- Each table contains a minimal amount of information and is related to other tables by a key
- For example, an Employee table and a Department table might be related by an EmployeeID
- Key values appear in each related table

The diagram illustrates four tables in a relational database and their relationships:

- Sales Table:** Columns include Sales Invoice #, Date, Salesperson, and Customer #. It contains 6 records.
- Customer Table:** Columns include Customer #, Customer Name, Street, City, and State. It contains 4 records.
- Sales-Inventory Table:** Columns include Sales Invoice #, Item #, and Quantity. It contains 11 records.
- Inventory Table:** Columns include Item #, Unit Price, and Description. It contains 6 records.

Relationships (indicated by red arrows):

- Sales to Customer:** A one-to-many relationship where the Customer # in the Sales table is a foreign key to the Customer # in the Customer table.
- Sales-Inventory to Sales:** A many-to-one relationship where the Sales Invoice # in the Sales-Inventory table is a foreign key to the Sales Invoice # in the Sales table.
- Sales-Inventory to Inventory:** A many-to-one relationship where the Item # in the Sales-Inventory table is a foreign key to the Item # in the Inventory table.

- A query language (SQL: Structure Query Language) is used to manipulate data in the RDB
- For the db in the example we might write this to find all its ordered by a particular customer:

```
select customer.id, customer.name, invoice.id,
sales.invoice, sales.item, inventory.item,
inventory.description
FROM customer, invoice, sales, inventory
WHERE sales.invoice = invoice.id AND sales.item =
inventory.item
```

- That earlier diagram is a form of ERD or entity-relationship diagram
- The ERD is a graphical way of showing how the tables are related
- Related in this case means that the tables share a common key
- Most DBs have tools to report ERDs, or you can start with an ERD and build tables from it

## 13. Hitting a mysql db from an app in PHP

- In the app, a connection is made to the db, queries are passed to the db, and results are stored in variables for processing

```
$db = mysql->connect('localhost:3307', 'accountDB', 'passwd');
$query='select * from accounts WHERE uid =' + uid;
$result = $db->query($query);
while ($row = $result->fetchRow()) {
    echo $row;
```

- }

#### 14. Relational db plusses and minuses

- a. RDBs tend to be very good at structured, related data
- b. We usually normalize the data so that each table only contains a very specific piece of information; relationships are used to tie these pieces together
- c. For example, a table of cities and zip codes would be used to pull data into a row that had a user's city...the user row would just have an id into the city table
- d. The minus is that complex relationships can take significant resources (time, memory) to build when needed
- e. We try to mitigate this with as much memory for caches and indexes as we can afford

#### 15. Views

- a. It's common to have a relational set of data that's used across several apps
- b. For example, we might want a user's name, address, city, and list of accounts
- c. These bits might come from several related tables, and it would be time-consuming to have to construct the query in either the app or the DB engine
- d. Instead we create a view, which is essentially a pre-built query that we can use instead of the full complex query
- e. The view ends up behaving like its own table...we do the lookup once and the db stores the result (not the same as caching)
- f. Views sometimes span multiple data sources, but usually they are constrained to a single DB
  - i. When we need a view that spans DBs we use middleware instead
- g. The view is an abstraction of the query
- h. In most DBs the views are pre-compiled for optimization
- i. We might also want a series of steps to happen when a query is made...these are called stored procedures or just stored procs

#### 16. Prepared statements

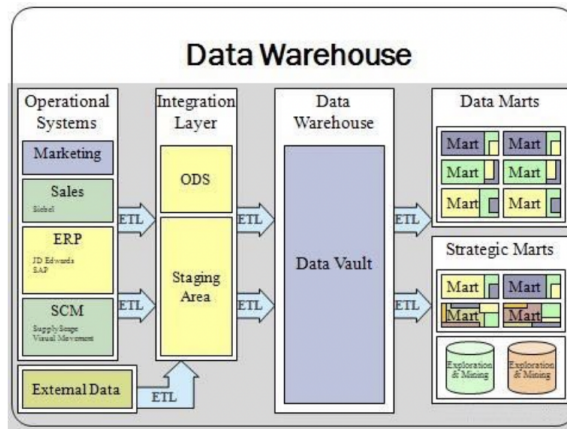
- a. This is a template that is parameterized and precompiled on the db side
- b. For example  
insert into contacts (name, phone) values (?, ?);
- c. This is compiled once and validated
- d. Normal statements are compiled and processed each time
- e. You can see significant performance enhancements

#### 17. The architect vs the DB guru

- a. One point of conflict in design is: Where does related data get assembled?
- b. The DB guru will want to do it in-db because it is faster...the db is optimized for that sort of thing
- c. The architect will want to do it in the app, because it can be optimized there
- d. Who is right?

## 18. Data warehouses

- a. There's a large family of apps that have heavy read-only requirements for data...reporting is a big example
- b. The data itself might be changing periodically, but for reporting purposes we only need most-recent data
- c. A common pattern for this problem is a data warehouse
- d. The warehouse collects and combines data from several sources into views
- e. The views are pushed into function-specific data marts for consumption by apps
- f. In most cases these views are denormalized...they contain one table per view



- g. These data marts / warehouses are often coupled with a reporting language or app
- i. Crystal Reports, for example, is a popular one
- j. The apps allow non-db users to build reports using drag-and-drop
- k. The reports are compiled into what is essentially a stored procedure in db and executed against current data

## 19. Non-relational DBs

- a. For web-based apps the trend has been toward non-relational DBs such as MongoDB
- b. The DBs are document-based rather than table based
- c. The term noSQL is used for this style...it doesn't necessarily mean 'no SQL' but rather 'non-relational'
- d. noSQL DBs trade simplicity for performance
- e. There are quite a few of them ...
- f. The idea behind non-relational DBs is to store all of the data necessary for a record into a single object
- g. Relationships are folded into that single object
- h. In the sales example from earlier, the invoices and items would be folded into a single customer object with arrays for invoices
- i. There can potentially be a lot of duplication
  - i. For example, the description of an item, or the address of a customer, might be repeated in hundreds of documents

- j. Turns out the answer is 'not all the time'
- k. In typical web-based applications the focus is on a user and his or her interaction with a site
- l. For this use case, a fully denormalized data store (ala noSQL) can perform better than a fully relational (highly normalized) DB
- m. For the corner cases where we really do need relational data, we take the performance hit

## 20. ORMs, ORDs, and impedance

- a. In most case an RDB does not map one-to-one onto the objects in the application code
- b. Work is required to move data from app to DB; the pattern we use is an ORM (object-relational mapping), basically a way to take RDB table info and store it in an object
- c. The problem is that the data types in the DB don't necessarily match the way we want to use them in the app...this is called an impedance mismatch
- d. In document-based DBs (such as mongoDB), there are few or no restrictions on data types in a record, reducing or elimination impedance
- e. However, this means that there are basically no filters on what goes into the mongoDB document store...the risk is GIGO
- f. Libraries such as Mongoose provide ODM (object-document mapping) services on top of mongoDB
- g. ODM is the equivalent of ORM but for non-relational DBs
- h. Bottom line is that ODM on top of document-based DBs provide a cleaner, more performant way to handle data in most applications

## 21. A note about ODM

- a. It's certainly possible to throw random, possibly unrelated JSON at a document-based DB like mongoDB, and then go and operate on it
- b. However, mongoDB does essentially no type checking, and the risk is that you'll be working on undefined data
- c. Using ODM tools such as Mongoose allow you to place schema-based constraints on types and ensures that data in the store is uniform

## 22. Using noSQL DBs

- a. We'll use mongoDB and Javascript for our examples
- b. In addition, Mongoose will provide ODM services through schemas
- c. The TL;DL is that we map Javascript objects 1:1 to database documents
- d. Mapping is done through JSON

## 23. JSON

- a. JSON = Javascript Object Notation
- b. It's a human-readable way to store and transmit data that maps directly to an object

- c. JSON has pretty much replaced XML in web apps, though there are still a huge number of AJAX based apps running on the web
- d. NB: You might see references to AJAX (asynch Javascript and JSON)
- e. JSON is the basic format for noSQL databases such as mongoDB

```
{
  "message": "Successfully retrieved all TODOs.",
  "todos": [
    {
      "_id": "56cf81a60e444bc7faa2ed43",
      "title": "Tonight's TODO!",
      "description": "Some description or not",
      "details": "this is not a detail",
      "_v": 0
    },
    {
      "_id": "56cf836f0e444bc7faa2ed44",
      "title": "new title",
      "description": "sdfs",
      "details": "hhhh junk",
      "_v": 0
    },
    {
      "_id": "56ce309851068982a09a573e",
      "title": "New todo today!",
      "description": "Some or not",
      "details": "Details now!",
      "_v": 0
    }
  ]
}
```

### Mapping between JS model and JSON mongoDB document

```
var mongoose = require('mongoose');
var Schema = mongoose.Schema;

var todoSchema = new Schema({
  title: String,
  description: String,
  details: String,
  owner: String
});
var Todo = mongoose.model('Todo', todoSchema);
```

f. }

## 24. Data directory

- a. By default mongoDB looks for data in /data/db
- b. It'll exit at start if the directory doesn't exist
- c. You also can point to any directory at startup, i.e.
  - i. mongod& —dbpath ~/data
  - ii. Might have to fiddle with permissions depending on path

## 25. Starting mongoDB

- a. The database runs as a daemon and listens for connection on port 27017 by default
  - i. This can be changed by editing resource file
  - ii. You might need to hunt for mongod.conf...it'll end up in the install directory
  - iii. Useful to create an alias in /etc to make it easy to find the second time (via ln -s)
- b. Startup
  - i. mongod&
  - ii. mongod -dbpath ~/data& (if you've moved the data directory to ~)

## 26. Security

- a. The base installation has authentication TURNED OFF by default
- b. This means anyone with access to the port has access to all databases
- c. (This is why there was a huge mongoDB ransomware attack a few months ago)
- d. Authentication is set up in mongod.conf
  - i. Requires definition of users and roles



- ii. Docs are at <https://docs.mongodb.com/manual/reference/configuration-options/#security-options>
- iii. Lots of info on the web, too, especially after the attacks
- e. In class we'll use defaults, but if you deploy to something other than your local machine you'll want to set up authentication

## 27. Import sample collections from mongoDB

- a. Grab the sample data
  - i. `curl https://raw.githubusercontent.com/mongodb/docs-assets/primer-dataset/primer-dataset.json > mongo-sample-collection.json`
- b. Import into running mongoDB instance
  - i. `mongoimport --db test --collection restaurants -- drop --file ./mongo-sample-collection.json`

## 28. Model and View development

- a. Typical workflow:
  - i. Business decides what data should be show on each screen / state
  - ii. Designers create HTML frames / mockup (tested FE)
  - iii. Data / back-end developers provide views (tested BE)
  - iv. Views are wired into FE via model/view

## 29. Caches

- a. API network calls are extremely expensive (in terms of processing time)
- b. Caching provides a 'close' store for frequently used data
- c. The risk is that the data in cache is older than what is in the database
- d. There are two general ways to approach this (ignoring client-side caches)

## 30. Caching strategies

- a. A write-through cache stores data in cache and immediately stores it in the DB; confirmation is delayed until the write is complete
- b. A write-back cache stores write data in cache and immediately confirms the write, then updates the DB when it becomes necessary
- c. The big downside of caches is coherency

## 31. Cache coherency

- a. When you need a piece of data and it is in cache (a "cache hit", how do you know it is the most current value? Especially if there's more than one user/ thread updating the DB
- b. One strategy is to do a fast metadata query on the backend to check if the data has been updated
- c. Another strategy is to turn the radio up louder and hope for the best (a surprisingly common approach)

- d. Sometimes it really doesn't matter...if you have data that doesn't change much, just load it up into memory when the app starts (things like state abbrevs/names and so on)
- e. Most DBs will do a fair amount of caching on their own, but they tend not to be as aggressive as an app-side cache

### 32. HTTP is stateless

- a. Infrastructure data includes session-related items that help us give the user an end-to-end experience
- b. HTTP is not naturally state-aware...each click (and in fact each element on a web page) is treated as a completely new transaction
- c. We use data both on the back end and stored in local cookies to give the user the illusion of a session
- d. Each time a user clicks a site link, several data transactions might be necessary to set up the resulting page
- e. A common approach is to generate and use a session ID which is transmitted to the web server on each click, and is used to store data in a session table

### 33. Session IDs

- a. Most frameworks provide a simple way to generate, store, and transmit session tokens
  - i. In cookies, transmitted as part of the HTTP request
  - ii. As a parameter in the query string (either a POST or a GET), generated and embedded in the link
- b. Cookies are the most common, though there's always some subset of users who refuse to enable them

### 34. PHP

- a. Stands for PHP Hypertext Processor
- b. Cross-platform
- c. Server-side language
  - i. PHP is processed by web server plugin which produces HTML
  - ii. Requires overhead to process
- d. PHP is an embedded language...it is mixed in with HTML
- e. The tokens surround PHP code
- f. The language itself is similar to PERL and C++
- g. It's optimized for web pages...deep database features, HTML code generation and manipulation
- h. Many popular platforms, such as Drupal, are based on PHP

## 35. Getting data using PHP

```
<?php
$servername = "localhost";
$username = "username";
$password = "password";
$dbname = "myDB";

// Create connection
$conn = new mysqli($servername, $username, $password, $dbname);
// Check connection
if ($conn->connect_error) {
    die("Connection failed: " . $conn->connect_error);
}

$sql = "SELECT id, firstname, lastname FROM MyGuests";
$result = $conn->query($sql);

if ($result->num_rows > 0) {
    // output data of each row
    while($row = $result->fetch_assoc()) {
        echo "id: " . $row["id"]. " - Name: " . $row["firstname"]. " " . $row["lastname"]. "<br>";
    }
} else {
    echo "0 results";
}
$conn->close();

//Source: W3Schools
?>
```

a.

## 36. Sessions

- Session tables can help provide the appearance of a session
- Use session IDs to store things like last page visited, shopping cart info, preferences, and so on
- Normally there is one session row per logged in user, and the session data itself can either be stored in that row in a denormalized way, or in a relational way

## 37. Session variables (PHP)

<pre><b>PAGE 1</b> &lt;?php // Start the session session_start(); ?&gt; &lt;!DOCTYPE html&gt; &lt;html&gt; &lt;body&gt;  &lt;?php // Set session variables \$_SESSION["favcolor"] = "green"; \$_SESSION["favanimal"] = "cat"; echo "Session variables are set."; ?&gt;  &lt;/body&gt; &lt;/html&gt;</pre>	<pre><b>PAGE2</b> &lt;?php session_start(); ?&gt; &lt;!DOCTYPE html&gt; &lt;html&gt; &lt;body&gt;  &lt;?php // Echo session variables that were set on previous page echo "Favorite color is " . \$_SESSION["favcolor"] . "&lt;br&gt;"; echo "Favorite animal is " . \$_SESSION["favanimal"] . "."; ?&gt;  &lt;/body&gt; &lt;/html&gt;</pre>
---	--

a.