

## Policy Learning V: Reinforcement Learning IV

### 1. Policy Search

- a. Core idea:
  - i. Modify the policy until performance stops improving
- b. What kind of policies?
  - i. Parameterized:
    - 1. Only useful if policy uses (significantly) less memory than the # of states
- c. For example:

$$\pi(s) = \underset{a}{\operatorname{argmax}} \hat{Q}_{\theta}(s, a)$$

- d. When policy relies solely on Q-functions:
  - i. Policy Search  $\rightarrow$  learn Q functions
  - ii. Different than Q-learning
    - 1. Q-learning  $\rightarrow$  goal is to learn Q-function that is “close enough” to  $Q^*$
    - 2. policy search  $\rightarrow$  goal is to find Q-functions that are “good enough”
    - 3. These goals might not produce the same thing!
- e. Problem:
  - i. Policy is discontinuous (i.e. not a smooth function)
- f. Why is discontinuity a problem?
  - i. Some small change in  $\theta$  can cause drastic changes in policy
  - ii. Bad for gradients!
- g. Idea: combine discontinuous estimates into a continuous Prob:

$$\pi(s, a) = \frac{e^{\hat{Q}_{\theta}(s, a)}}{\sum_{a'} e^{\hat{Q}_{\theta}(s, a')}} \quad \left| \right.$$

- h. This is called the softmax function
  - i. A form of normalizing a vector  $\rightarrow$  produces probabilities
- i. Why softmax?
  - i. Almost deterministic when one action is vastly better than others
  - ii. Differentiable!

- j. One other nice property:

$$\begin{aligned} & \log \left( \frac{e^{\hat{Q}_\theta(s, a)}}{\sum_{a'} e^{\hat{Q}_\theta(s, a')}} \right) \\ &= \log \left( \frac{e^{\hat{Q}_\theta(s, a)}}{C} \right) \\ &= \hat{Q}_\theta(s, a) - \log(C) \end{aligned}$$

- k. Why is this important?

- i. Remember, we want policy to optimize

$$\mathbb{E}[R(\tau)] = \mathbb{E} \left[ \sum_t \gamma^t R(s_t) \right]$$

- ii. How can we represent  $\mathbb{E}$ ?

$$\mathbb{E}[f(x)] = \sum_x \text{Pr}[x] f(x)$$

- iii. Looks kinda similar!

$$\mathbb{E} \left[ \sum_t \gamma^t R(s_t) \right] = \sum_t \pi_\theta(s_t) \gamma^t R(s_t)$$

- l. We need to differentiate

$$\begin{aligned} & \frac{\partial}{\partial \theta_i} \mathbb{E} \left[ \sum_t \gamma^t R(s_t) \right] = \frac{\partial}{\partial \theta_i} \sum_t \pi_\theta(s_t) \gamma^t R(s_t) \\ &= \sum_t \frac{\partial}{\partial \theta_i} \pi_\theta(s_t) \gamma^t R(s_t) \\ &= \sum_t \gamma^t R(s_t) \frac{\partial}{\partial \theta_i} \pi_\theta(s_t) \\ &= \sum_t \gamma^t R(s_t) \pi_\theta(s_t) \frac{\frac{\partial}{\partial \theta_i} \pi_\theta(s_t)}{\pi_\theta(s_t)} \\ &= \sum_t \gamma^t R(s_t) \pi_\theta(s_t) \frac{\partial \log(\pi_\theta(s_t))}{\partial \theta_i} \\ &= \mathbb{E} \left[ \sum_t \gamma^t R(s_t) \frac{\partial}{\partial \theta_i} \log(\pi_\theta(s_t)) \right] \end{aligned}$$

m. Woah!

$$\nabla_{\theta} \mathbb{E} \left[ \sum_t \gamma^t R(s_t) \right] = \mathbb{E} \left[ \sum_t \gamma^t R(s_t) \nabla_{\theta} \log(\pi_{\theta}(s_t)) \right]$$

- n. Our gradient = gradient of policy scaled by how good the choices were
- o. Can take sample average to approximate expectation!
- p. Turns out there is a little more we can do here!

$$\nabla_{\theta} \mathbb{E} \left[ \sum_t \gamma^t R(s_t) \right] = \mathbb{E} \left[ \sum_t G_t \nabla_{\theta} \log(\pi_{\theta}(s_t)) \right]$$

- i.  $G_t$  = value of trajectory from that point on

## 2. Policy Search with REINFORCE

- a. Transforms search into a Monte-Carlo procedure
  - i. Samples one trajectory (i.e. plays one game)
  - ii. Records trajectory
- b. Updates policy w policy gradient (offline i.e. in between games)

### function REINFORCE

Initialise  $\theta$  arbitrarily

**for** each episode  $\{s_1, a_1, r_2, \dots, s_{T-1}, a_{T-1}, r_T\} \sim \pi_{\theta}$  **do**

**for**  $t = 1$  to  $T - 1$  **do**

$\theta \leftarrow \theta + \alpha \nabla_{\theta} \log \pi_{\theta}(s_t, a_t) v_t$

**end for**

**end for**

**return**  $\theta$

**end function**

- c. Repeat!

## 3. Can We do Better?

- a. Policy Gradient:

$$\mathbb{E} \left[ \sum_t G_t \nabla_{\theta} \log(\pi_{\theta}(s_t)) \right]$$

- b. What if  $G_t$  was its own function-approx  $G_{\theta}$  (could be  $Q_{\theta}(s_t, a_t)$ )
- c. Already have  $\pi_{\theta}$

## 4. Actor-Critic RL

- a. The “critic” = a function approximation for the “value” function

- i. Could be Q-value, utility value
- b. The “actor” = a function approximation for the policy
- c. Critic still obeys bellman equation, can use TD learning

$$\theta'_i \leftarrow \theta'_i + \alpha \left( R(s) + \gamma \hat{Q}_{\theta'}(s', a') - \hat{Q}_{\theta'}(s, a) \right) \frac{\partial \hat{Q}_{\theta'}(s, a)}{\partial \theta'_i}$$

- d. Actor gradients we just derived!

$$\theta_j \leftarrow \theta_j + \eta G_t \nabla_{\theta} \log \left( \pi_{\theta}(s_t) \right)$$