

## Policy Learning V: Reinforcement Learning V

## 1. The Advantage Function (A2C)

a. What happened to using the Bellman equation?

b. We can include it for  $G_{\theta'}$ i.  $G_{\theta'}$  just estimates the “value” of the current choice

ii. But we know choices are related!

1. How good is a choice?

$$2. \quad A(s_t, a_t) = Q_{\theta'}(s_t, a_t) - V_{\theta''}(s_t)$$

iii. This is called the advantage function

1. Do we need two function-approximations, one for  $Q$  and one for  $V$ ?

2. No!

$$\text{iv.} \quad Q(s_t, a_t) = \mathbb{E} \left[ R(s_{t+1}) + \gamma V(s_{t+1}) \right]$$

v. Therefore

$$A(s_t, a_t) = R(s_{t+1}) + \gamma V_{\theta''}(s_{t+1}) - V_{\theta''}(s_t)$$

c. We only need  $V_{\theta''}$  and  $\pi_{\theta}$ !

## 2. Offline Actor-Critic RL

a. Could also do this offline

i. Sample  $N$  trajectories (i.e. play  $N$  games) and record trajectories

ii. Build supervised learning datasets


iii. For  $k$  training iterations:1. Update  $\pi_{\theta}$  and  $G_{\theta'}$ 

iv. Repeat

b. Useful for scaling:

i. A3C (Async A2C) = play lots of games in parallel (in explore step), update after

## 3. NN Q-Function



- a.
- b. NN takes place of tabular Q-function
- c. How to apply updates?
  - i. If we had a table:

$$Q(s, a) \leftarrow Q(s, a) + \alpha \left( R(s) + \gamma \max_{a'} Q(s', a') - Q(s, a) \right)$$

- ii. Since we have a NN...cannot adjust Q value directly (like in above equation):

$$\theta_i \leftarrow \theta_i + \alpha \left( R(s) + \gamma \max_{a'} \hat{Q}_{\theta}(s', a') - \hat{Q}_{\theta}(s, a) \right) \frac{\partial \hat{Q}_{\theta}(s, a)}{\partial \theta_i}$$

- d. Gradient descent!
  - i. Error (loss function) =

$$R(s) + \gamma \max_{a'} \hat{Q}_{\theta}(s', a') - \hat{Q}_{\theta}(s, a)$$

#### 4. Training a NN Q-function

- a. For now, let us assume exploration function f is randomized impl.
  - i. With some probability p, ignore policy and choose random action
  - ii. With remaining probability 1-p, follow policy
- b. Let us observe the transition (s, a, s') where we have recorded R(s) and R(s')
  - i. Fix the policy (i.e. just use the NN):
    1. Calculate TD error by iterating over possible actions a' in s'

$$e = R(s) + \gamma \max_{a'} \hat{Q}_{\theta}(s', a') - \hat{Q}_{\theta}(s, a)$$

- ii. Backprop e through the network treating e like it is any other error function