



Deep Learning: Artificial Neural Networks: Design and Implementation; Workflow for Classification Tasks; Evaluating Classifiers

1. Introduction to Deep Learning

- a. Deep Learning refers to Supervised Learning using an Artificial Neural Network, which has the following features:
- b. It is a network/graph of small computation units called artificial neurons, loosely modeled on the neurons in our brains, which send signals to each other. The signals are floating-point numbers.
- c. The network is typically organized in layers: the first layer is the input layer, the last is the output layer, and others are called hidden layers.
- d. The input layer is array/vector of floats, and the output layer produces an array of floats. Thus, the network computes a function from vectors to vectors



e.




f.


g. Each neuron:

- i. Is connected to each neuron in the previous layer, and each connection has a weight or parameter which determines the strength of the signal (importance of this input to the neuron);
- ii. Performs logistic regression, using the sigmoid or other non-linear activation function.

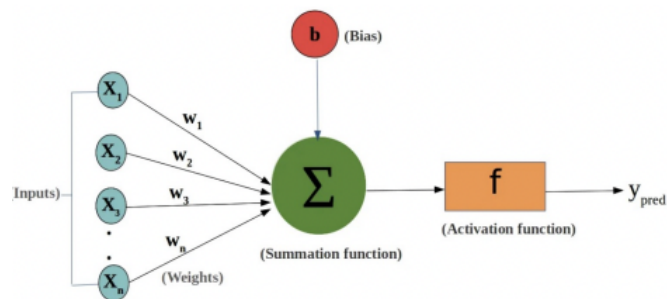
- h. Gradient descent is used to learn the weights to minimize some cost function on the outputs.



i.



j.




k.

- l. Features of artificial neural networks:

- i. Additional layers may perform data aggregation (e.g., convolution and pooling), dropout (forgets some of the weights), or other kinds of data manipulation (e.g., softmax = transforming the output into a probability distribution).
- ii. In a feedforward network, the network transforms an array of floats through the layers into another array of floats; in a sequence model, the inputs and outputs are sequences of vectors; and recurrent layers have cyclical connections which act as memory.
- iii. BERT, GPT, and other large networks learn to pay attention to complex patterns in the input sequence (e.g., words in a sentence).

A neuron is a higher-dimensional version of our logistic regression algorithm:


$$\sigma(m * x + b)$$



m.


n. One possible activation function f is the sigmoid from logistic regression:

$$y = \sigma(z) = \frac{1}{1 + e^{-z}}$$




i.

o. Both non-linear activation functions besides sigmoid are more often used



Hyperbolic Tangent (Tanh)



Rectified Linear Unit (Relu)


i.

ii. Hyperbolic tangent works better than sigmoid in recurrent calls

p. A small amount of Linear Algebra can be used to compactly specify the logistic regression performed by a neuron.

q. Inputs and weights/parameters are just vectors:

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_m \end{bmatrix} \quad \mathbf{w} = \begin{bmatrix} w_1 \\ w_2 \\ \vdots \\ w_m \end{bmatrix}$$



i.

r. Often written as tuples


i. $\mathbf{X} = (x_1, x_2, \dots, x_m).$ $\mathbf{W} = (w_1, w_2, \dots, w_m).$

s. Then the neuron performs logistic regression by performing a dot product of inputs \mathbf{x} and weights \mathbf{w} ,

i. $\mathbf{W} \cdot \mathbf{X} = (w_1 \cdot x_1) + \dots + (w_k \cdot x_k) = \mathbf{W}^T \mathbf{X}$

t. and then applying the activation function f :


$$\sigma(m * x + b) \qquad a = f(\mathbf{W}^T \mathbf{X}).$$



i. $\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_m \end{bmatrix} \qquad \mathbf{w} = \begin{bmatrix} w_1 \\ w_2 \\ \vdots \\ w_m \end{bmatrix}$

u. A small but important detail: Each neuron has a bias term, because they are simply doing logistic regression! (each neuron will have a bias)

$$\sigma(m * x + b * 1.0)$$




i. \rightarrow you always have a '1.0' to match 'b' in matrix multiplication

v. This connection has a weight (the value b) but is not connected to the inputs; it serves to scale the output, just as b does in linear regression.


w. A convenient way to encode this is to assume that all input vectors to all neurons contain a constant 1.0 value:

i. $\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_{m-1} \\ 1.0 \end{bmatrix} \qquad \mathbf{w} = \begin{bmatrix} w_1 \\ w_2 \\ \vdots \\ w_{m-1} \\ b \end{bmatrix}$



- x. The simplest network is simply a row of such neurons, where
 - i. Each has its own weight/parameter vector, but
 - ii. Shares the same input vector

Activations (outputs)



Inputs

- y.
 - i. They can work together to learn different thing about the inputs
 - ii. Each will do logistic regression using its weights on the input and output an activation (each separately)
 - iii. Vector-to-vector function
 - iv. It works simultaneously (everything in one layer in principle)
- z. The inputs and outputs for a layer are vectors

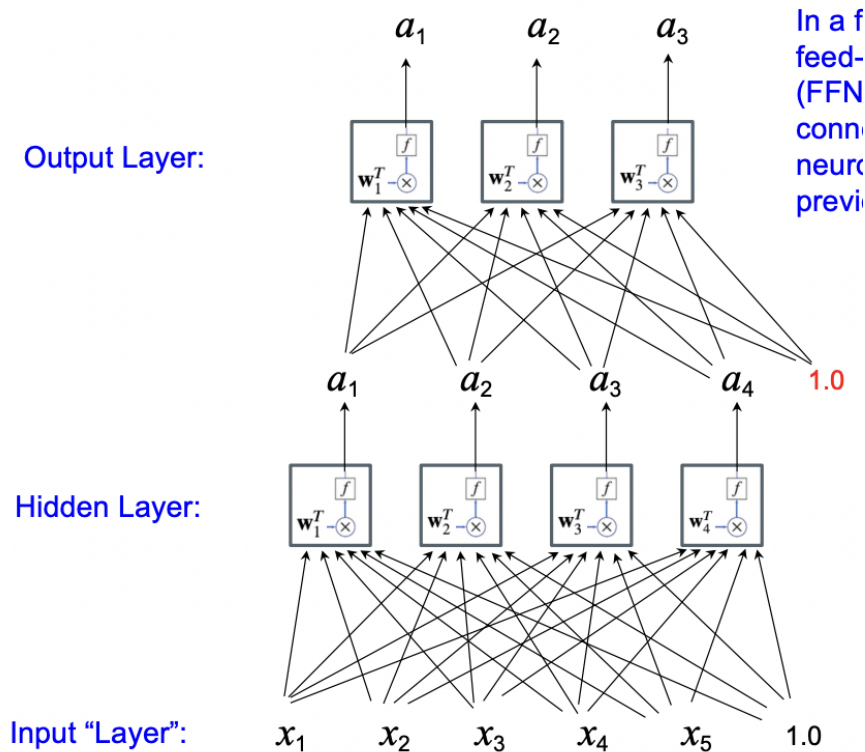
- i. $\mathbf{a} = (a_1, a_2, a_3, a_4)$

- aa. The weights/parameters for a layer form a matrix

$$\mathbf{W} \mathbf{x} = \begin{bmatrix} w_{1,1} & w_{1,2} & w_{1,3} & w_{1,4} & w_{1,5} & b_1 \\ w_{2,1} & w_{2,2} & w_{2,3} & w_{2,4} & w_{2,5} & b_2 \\ w_{3,1} & w_{3,2} & w_{3,3} & w_{3,4} & w_{3,5} & b_3 \\ w_{4,1} & w_{4,2} & w_{4,3} & w_{4,4} & w_{4,5} & b_4 \end{bmatrix} \times \begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ x_4 \\ x_5 \\ 1.0 \end{bmatrix} = \begin{bmatrix} \mathbf{w}_1^T \mathbf{x} \\ \mathbf{w}_2^T \mathbf{x} \\ \vdots \\ \mathbf{w}_n^T \mathbf{x} \end{bmatrix}$$

$$\mathbf{x} = (x_1, x_2, x_3, x_4, x_5, 1.0)$$

- i.
 - ii. The ws are weights that are learned



bb.

- i. In a fully connected feed-forward network (FFNN), each layer is connected to each neuron or input in the previous layer
- ii. Notice that each neuron has a bias weight
- iii. Each one of them is doing classification but they are working together
- iv. The bigger the network, there are problems introduced
 1. Cost of multiplication
 2. Searching in the highly dimensional space
 3. Vanishing weights/parameters \rightarrow the probabilities are so small that they disappear/vanish (Same thing can happen here)
 4. Some numbers can get very small or big (happens on deeper networks but not usually in wider networks)


cc. Again, a little bit of linear algebra can make this a whole lot simpler:

$$W \mathbf{x} = \begin{bmatrix} w_{1,1} & w_{1,2} & w_{1,3} & w_{1,4} & w_{1,5} & b_1 \\ w_{2,1} & w_{2,2} & w_{2,3} & w_{2,4} & w_{2,5} & b_2 \\ w_{3,1} & w_{3,2} & w_{3,3} & w_{3,4} & w_{3,5} & b_3 \\ w_{4,1} & w_{4,2} & w_{4,3} & w_{4,4} & w_{4,5} & b_4 \end{bmatrix} \times \begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ x_4 \\ x_5 \\ 1.0 \end{bmatrix} = \begin{bmatrix} \mathbf{w}_1^T \mathbf{x} \\ \mathbf{w}_2^T \mathbf{x} \\ \vdots \\ \mathbf{w}_n^T \mathbf{x} \end{bmatrix}$$

i.

Activations (outputs)

$$\mathbf{a} = (a_1, a_2, a_3, a_4)$$




$$\mathbf{a} = f(W\mathbf{x})$$

$$\mathbf{x} = (x_1, x_2, x_3, x_4, x_5, 1.0)$$

ii.


Inputs

dd. Our FFNN:




i.

$$\mathbf{a}^2 = (a_1^2, a_2^2, a_3^2)$$



$$\mathbf{a}^1 = (a_1^1, a_2^1, a_3^1, a_4^1) + [1.0]$$



$$\mathbf{x} = (x_1, x_2, x_3, x_4, x_5, 1.0)$$

ii.

$$\mathbf{a}^2 = f(W_2(f(W_1\mathbf{x}) + [1.0]))$$

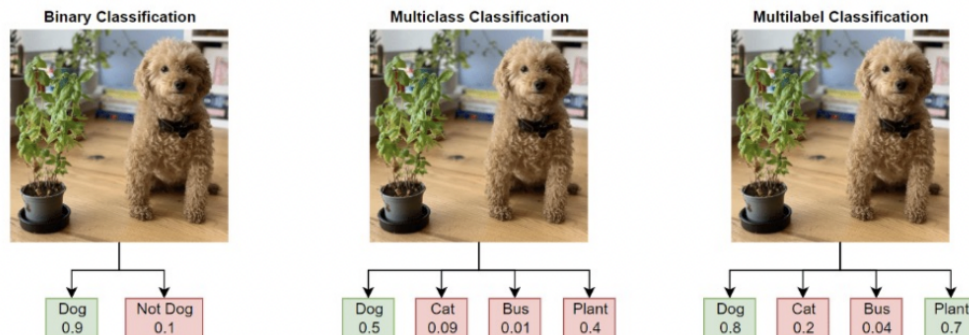
1. Do non-linear activation function every layer

2. Classification Methods: Supervised ML

- a. Input:
 - i. a fixed set of classes $C = \{c_1, c_2, \dots, c_J\}$
 - ii. a randomly-permuted set of labeled documents $(d_1, c_1), \dots, (d_n, c_n)$ split into
 1. a training set $(d_1, c_1), \dots, (d_m, c_m)$
 2. a testing set d_{m+1}, \dots, d_n (labels withheld)
- b. Output:
 - i. A classifier $\gamma : d \rightarrow c$ trained the training set
 - ii. The testing set with labels calculated by γ
 - iii. Test results (confusion matrix, metrics, etc.)

3. Classification: Binary, Multiclass and Multilabel

- a. In Binary Classification, we have 2 labels, and we must choose one; often this is phrased as “something” or “not something” (spam or ham, misinformation or not, etc.)
- b. In Multiclass Classification, we have more than 2 labels, and our task is to assign a single label to each sample:
- c. In Multilabel Classification, we have more than 2 labels, and our task is to assign any appropriate labels (not just one):



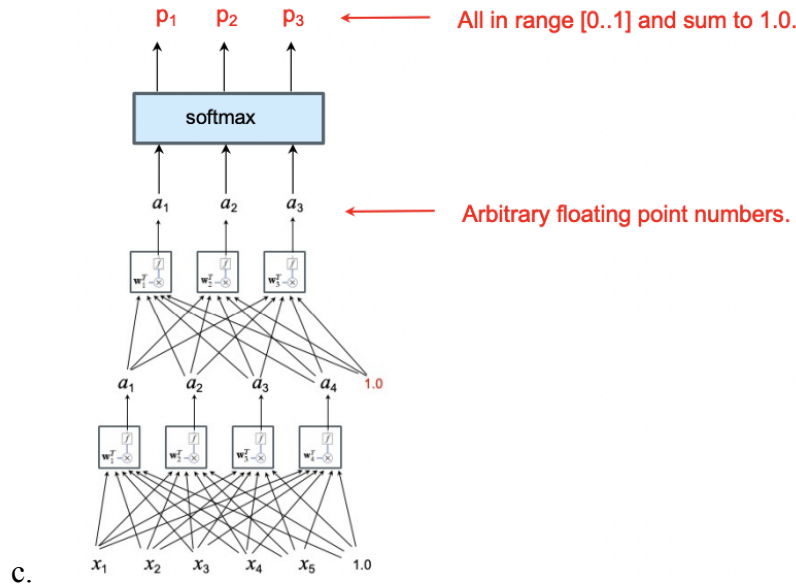
- d.
 - i. Multilabel has a threshold \rightarrow anything above the threshold is in the picture

4. Classification Methods: Supervised ML

- a. There are many different kinds of classifiers for labeled data
 - i. Naïve Bayes
 - ii. Logistic regression
 - iii. Neural networks

5. Classification with Deep learning

- a. All of these types of classification are typically implemented by having the network output a probability distribution on all the labels.
- b. To convert the output values into a distribution, we used a generalization of the sigmoid called the softmax



- i. No guarantee that the numbers coming out are good probability distributions (might not be between 0 and 1 and add to 1 and etc.)
- ii. Use softmax
- iii. Take any sequence of floating numbers and turn it into a probability distribution (it will maintain the order but scale the value)
- iv. Softmax exaggerates the differences

6. Introduction to Deep Learning

- a. Softmax = a generalization of sigmoid which scales k numbers into a probability distribution

$$s(z) = \frac{e^z}{e^z + 1}$$

- i.
- b. For a vector z of dimensionality k , the softmax is

$$\text{softmax}(z) = \left[\frac{\exp(z_1)}{\sum_{i=1}^k \exp(z_i)}, \frac{\exp(z_2)}{\sum_{i=1}^k \exp(z_i)}, \dots, \frac{\exp(z_k)}{\sum_{i=1}^k \exp(z_i)} \right]$$

- i.
- ii. Easy to manipulate, easy to take derivatives, and etc.
- c. Example

```
5 A = [1.2, 0.5, 3.2, 2.2, 1.9]
6 softmax(A)
array([0.07343397, 0.03646623, 0.54260772, 0.19961422, 0.14787785])
```

```
1 A = [1.2, 2.3]
2 softmax(A)
array([0.24973989, 0.75026011])
```

i.

1. 3.2 is not that far away from 2.2 but the softmax value is quite far away (emphasizing the larger ones)

7. Data Format for Deep Learning Classifier

- The training set consists of a matrix of features X and a vector of labels Y .
- For each input observation $x(i)$, we have a vector of features $[x_1, x_2, \dots, x_n]$. Feature j for input $x(i)$ is x_j , more precisely $x_j(i)$.

	x_1	x_2	x_3	...	x_{19}	x_{20}	y
$x^{(1)}$	0	1	0	0	0	1	0
$x^{(2)}$	1	0	1	1	0	0	1
$x^{(3)}$	0	0	0	0	1	3	0
$x^{(4)}$	2	0	0	0	0	1	0
$x^{(5)}$	0	0	0	0	0	0	0
$x^{(6)}$	0	0	0	0	1	0	0

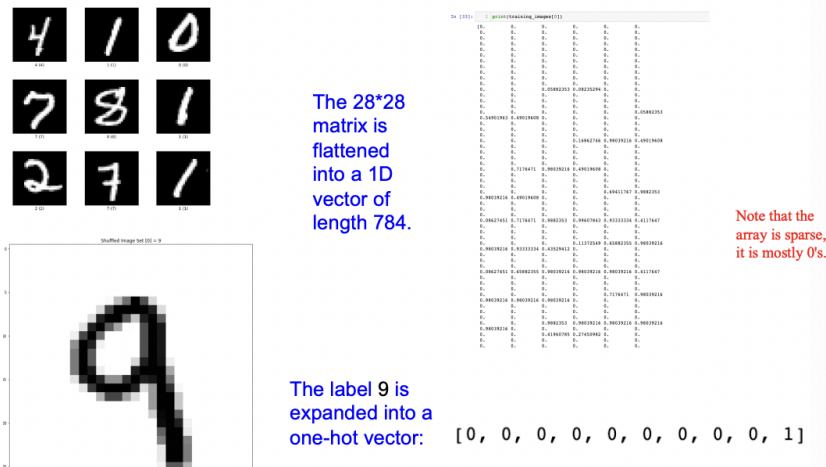
Feature Vectors X

Labels Y

c.

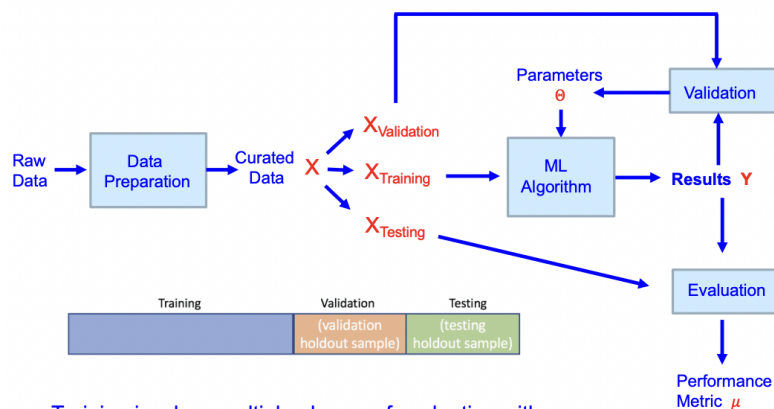
8. Multiclass Classification

- We will consider the MNIST database of handwritten digits: this is the “Hello World” of deep learning.
- The MNIST digit database consists of 70,000 28x28 BW pixel images, stored as $28 \times 28 = 784$ floating-point numbers in the range $[0..1]$; labels are integers 0..9:



c.

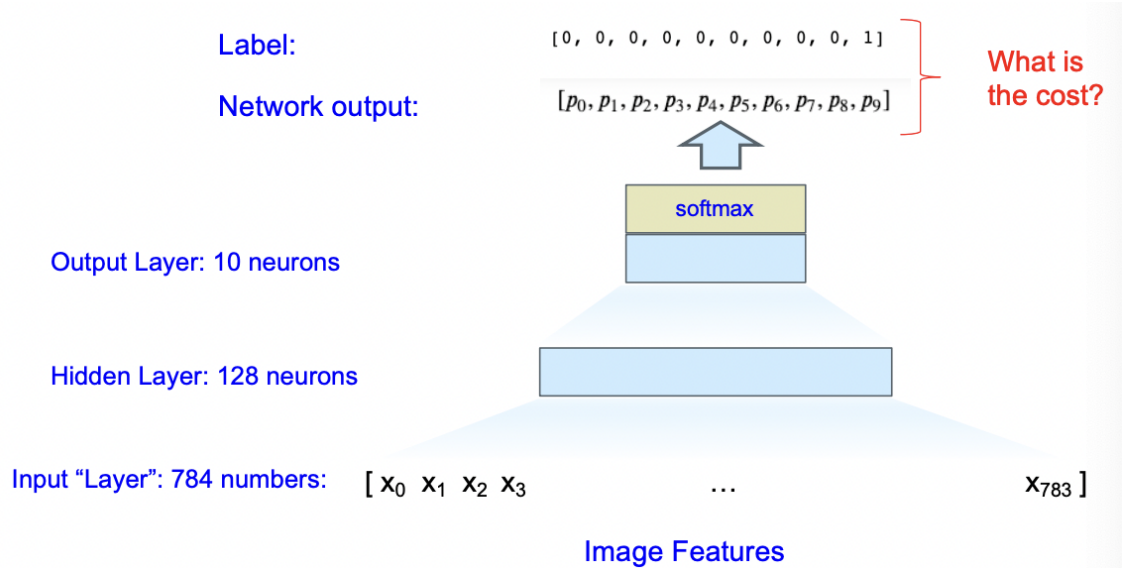
9. Recall: Supervised Machine Learning Workflow



a.

10. Cost/Loss Function for Classification

- a. One more detail! What is the cost (loss) function used with the output of a classifier?



- b.
- c. We need to compare two probability distributions

Label: $[0, 0, 0, 0, 0, 0, 0, 0, 0, 1]$ \mathbf{y}

Network output: $[p_0, p_1, p_2, p_3, p_4, p_5, p_6, p_7, p_8, p_9]$ $\hat{\mathbf{y}}$

to measure how different they are. The function used to do this is called the Cross-Entropy Loss:

$$\hat{\mathbf{y}} = \text{softmax}(\mathbf{a}) \quad CE(\hat{\mathbf{y}}, \mathbf{y}) = -\frac{1}{n} \sum_{i=1}^n y_i \log(\hat{y}_i)$$

- d. One great feature of this cost function is that the derivative of the softmax plus CE Loss function is absurdly simple: it is just the difference of the two distributions

$$\frac{\partial CE(\hat{y}_i, y_i)}{\partial z_i} = \hat{y}_i - y_i$$

i.